

Obesity prediction

The Kaggle link of this dataset is <https://www.kaggle.com/mrsimple07/obesity-prediction> and the usability rating is 10.

In this project we are trying to predict the Obesity category of a person based on their gender, age, height, weight, BMI and physical activity level. The target is a categorical variable with values normal weight, obese, overweight and underweight. An individual based on their features described earlier will be classified to fall under one of this target categories. The total records in the dataset are 1000.

Data Exploration

- **Data preparation issues:** As gender is categorical, this needs to be encoded to be numerical for comparison.

```
# Data Exploration
```

```
print(dataset.describe())
```

	Age	Height	Weight	BMI \
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	49.857000	170.052417	71.205769	24.888317
std	18.114267	10.309971	15.509849	6.193912
min	18.000000	136.115719	26.065730	8.470572
25%	35.000000	163.514205	61.129629	20.918068
50%	50.000000	169.801665	71.929072	24.698647
75%	66.000000	177.353596	81.133746	28.732132
max	79.000000	201.419670	118.907366	50.791898

	PhysicalActivityLevel
count	1000.000000
mean	2.534000
std	1.116284
min	1.000000
25%	2.000000
50%	3.000000
75%	4.000000
max	4.000000

Obesity Prediction with Random Forest

```
print(dataset.isnull().sum())
```


```
Age          0
Gender       0
Height       0
Weight       0
BMI          0
PhysicalActivityLevel  0
ObesityCategory  0
dtype: int64
```

```
print(dataset.info())
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   Age                        1000 non-null  int64
1   Gender                    1000 non-null  object
2   Height                    1000 non-null  float64
3   Weight                    1000 non-null  float64
4   BMI                      1000 non-null  float64
5   PhysicalActivityLevel     1000 non-null  int64
6   ObesityCategory          1000 non-null  object
dtypes: float64(3), int64(2), object(2)
memory usage: 46.9+ KB
None
```

Feature Analysis


Below are some summary statistics of the dataset. Firstly, we would like to understand if the dataset has a good representation of the population. To this end, we start by looking at the histograms of age and gender.



Obesity Prediction with Random Forest




Based on the gender, the data is equally distributed.


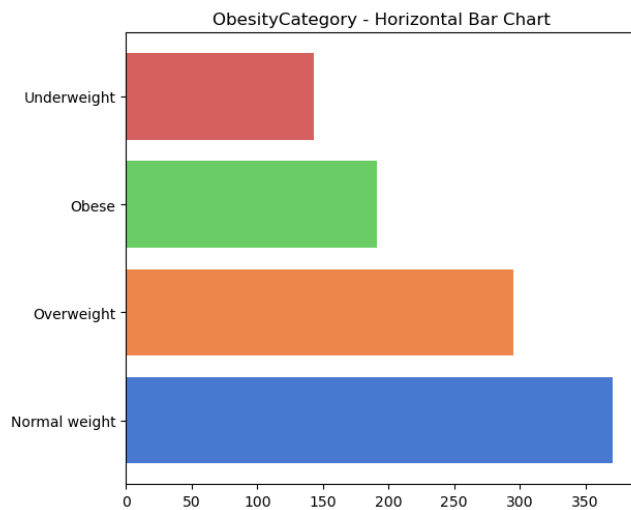


Obesity Prediction with Random Forest

Gender - Donut Chart

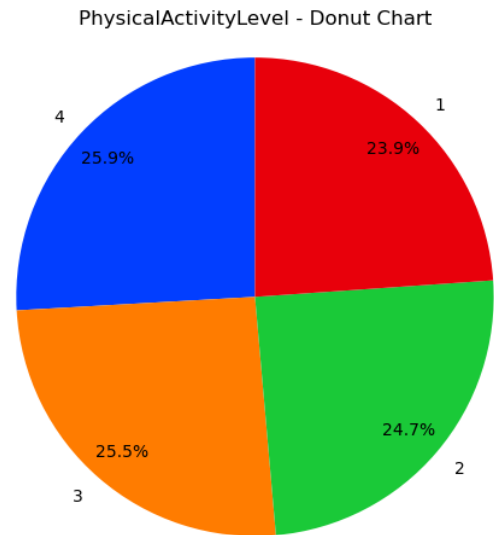



Plotting the number of people under each target label of obesity category.




Based on physical activity level,

Obesity Prediction with Random Forest



As BMI is based on height and weight, which is evident from the below correlation map. Hence those features can be dropped.



To be able to plot obesity category along with the other variables we convert obesity category into the following ordered integer values.

Obesity Prediction with Random Forest


Underweight = 0

Normal weight = 1

Overweight = 2

Obese = 3

We see the following by plotting obesity category vs. BMI.



Attribute selection

The relationship between BMI and features such as age, gender, height and weight are more interesting. Typically, BMI is calculated based on these features. The obesity category itself then is typically based on the range for the BMI. Therefore, if we know BMI of a person then we can in principle ignore all the other features to predict the obesity category. And this is evident from the below.

```
[30]: X_full = obesity_data.copy(deep=True)
      y_full = X_full.pop('ObesityCategoryInt')
      X_full['Gender'] = X_full['Gender'].apply(lambda g: 1 if g == 'Male' else 0)
      X_full.drop(columns=['ObesityCategory'], inplace=True)

      feature_selector = SequentialFeatureSelector(DecisionTreeClassifier(), direction='forward', tol=.001)
      feature_selector.fit(X_full, y_full)
      feature_selector.feature_names_in_[feature_selector.get_support()]

[30]: array(['BMI'], dtype=object)

[31]: feature_selector = SequentialFeatureSelector(DecisionTreeClassifier(), direction='backward', tol=-.001)
      feature_selector.fit(X_full, y_full)
      feature_selector.feature_names_in_[feature_selector.get_support()]

[31]: array(['BMI'], dtype=object)
```

Obesity Prediction with Random Forest


Our Model- Decision Tree with Grid Search cross validation

It appears that it is sufficient to find a model that maps BMI to Obesity Category. Given the strong split between categories, we could consider a decision tree. We indeed see a very strong performance even with this simple model.

```
X_train, X_test, y_train, y_test = train_test_split(obesity_data[['BMI']],
                                                    obesity_data['ObesityCategoryInt'],
                                                    test_size=0.2,
                                                    stratify=obesity_data['ObesityCategoryInt'],
                                                    random_state=121)

param_grid = {
    'max_depth': [2, 3, 4, 5, 6, 7, 8, 9, 10]
}

cross_validation = GridSearchCV(DecisionTreeClassifier(), param_grid, cv=StratifiedKFold(n_splits=5))
cross_validation.fit(X_train, y_train)
```



And we are getting 100% accuracy


```
print('Results on test set :', cross_validation.best_estimator_.score(X_test, y_test))
Results on test set : 1.0
```

We also did a comparison with KNN and Random Forest and got 100% accuracy score.

Comparison with different classifiers

The dataset was run through different classifiers and below the comparison chart of the accuracy score.

	Model	Accuracy
0	DecisionTreeClassifier	1.0
1	RandomForestClassifier	1.0
2	KNeighborsClassifier	1.0

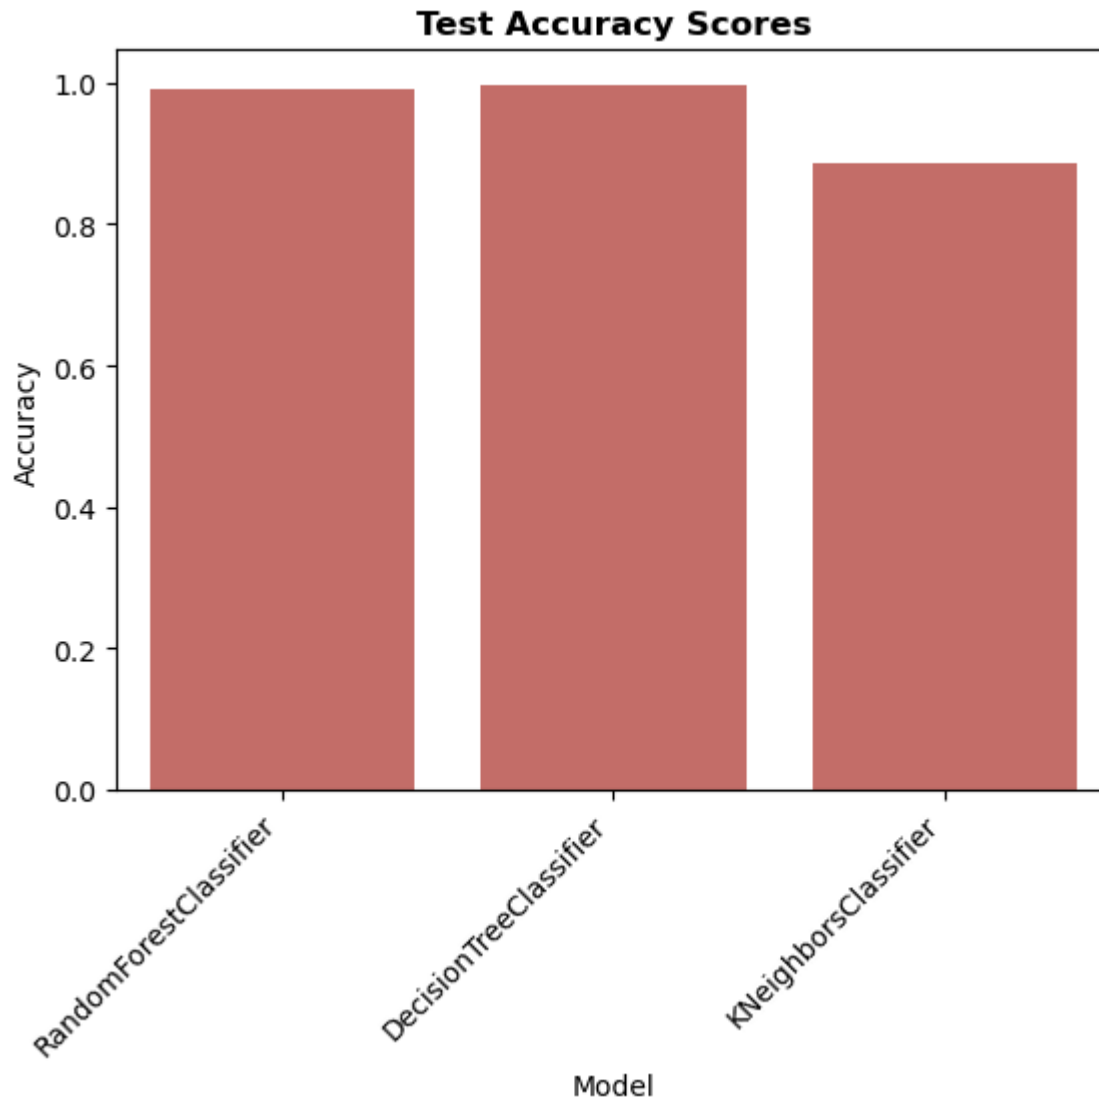


Contribution by other Kaggle participants

All the Kaggle participants seem to miss the fact that BMI alone is a sufficient predictor of obesity category. Most of the participants have used a mixture of different models to achieve an accuracy that is not any better than the model Decision Tree/Random Forest/KNN with Grid Search. We also tried to implement the 3 models using all features like other Kaggle contributors and obtained the below accuracy score.

Obesity Prediction with Random Forest

	Model	Accuracy
0	RandomForestClassifier	0.990
1	DecisionTreeClassifier	0.997
2	KNeighborsClassifier	0.887



Code implementation is present in the Github link – [Obesity Prediction](#)

Project Team – Sheema Selvam & Jayasree Jandhyala

Responsibilities of the team – Analyzing the Kaggle dataset, feature selection, comparing different Machine Learning models for classification

Sheema – Classification based on all features and comparison with different models with stratified k-fold validation

Obesity Prediction with Random Forest

Jayasree – Classification based on only BMI and comparison with different models with stratified k-fold validation and Grid Search